

Exam Task - Spam Message Classifier

December 19, 2018

1 SPAM message classifier

This task is based on the spam message dataset provided at <http://www.dt.fee.unicamp.br/~tiago/smsspamcollection/>. "The SMS Spam Collection is a set of SMS tagged messages that have been collected for SMS Spam research. It contains one set of SMS messages in English of 5,574 messages, tagged according to being ham (legitimate) or spam."

The data is being used to train a couple of our models as well check the accuracy of prediction. Then we this model will be integrated into our simple web application where you can input a text string get prediction results, if the message is indeed a spam.

1.1 1. Import and manipulate/check data

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

```
In [2]: # load data from csv file
data = pd.read_csv("./data/spam.csv", encoding='latin-1')
```

```
In [3]: data.shape
```

```
Out[3]: (5572, 5)
```

```
In [4]: data.head(5)
```

```
Out[4]:
```

	v1	v2	Unnamed: 2	\
0	ham	Go until jurong point, crazy.. Available only ...		NaN
1	ham	Ok lar... Joking wif u oni...		NaN
2	spam	Free entry in 2 a wkly comp to win FA Cup fina...		NaN
3	ham	U dun say so early hor... U c already then say...		NaN
4	ham	Nah I don't think he goes to usf, he lives aro...		NaN

	Unnamed: 3	Unnamed: 4
0	NaN	NaN
1	NaN	NaN
2	NaN	NaN

```

3      NaN      NaN
4      NaN      NaN

```

```

In [5]: # select and rename the columns relevant to us
data = data[['v2', 'v1']]
data = data.rename(columns={"v1": "label", "v2": "text"})

print(data.groupby('label').size())

```

```

label
ham      4825
spam      747
dtype: int64

```

```

In [6]: data.isnull().values.any()

```

```

Out[6]: False

```

```

In [7]: # convert the output data/ label to binary value
data['label'] = data.label.replace(to_replace=['ham', 'spam'], value=[0, 1])

```

```

In [8]: data.head(5)

```

```

Out[8]:
          text  label
0  Go until jurong point, crazy.. Available only ...      0
1              Ok lar... Joking wif u oni...      0
2  Free entry in 2 a wkly comp to win FA Cup fina...      1
3  U dun say so early hor... U c already then say...      0
4  Nah I don't think he goes to usf, he lives aro...      0

```

1.2 2. Visualization of data

We used a visualization example for wordcloud to get an overview of the words present in the spam and legitimate messages

1.2.1 2.1 Message length distribution per category

```

In [9]: plotdata = data
plotdata["message_length"] = data["text"].apply(len)

f, ax = plt.subplots(1, 2, figsize = (20, 6))


sns.distplot(plotdata[plotdata["label"] == 1]["message_length"], bins = 20, ax = ax[0])
ax[0].set_xlabel("Spam Message Word Length")

sns.distplot(plotdata[plotdata["label"] == 0]["message_length"], bins = 20, ax = ax[1])
ax[1].set_xlabel("Ham Message Word Length")

plt.show()

```

```
C:\ProgramData\Anaconda3\lib\site-packages\matplotlib\axes\_axes.py:6462: UserWarning: The 'normed' kwarg is deprecated, and has been "
warnings.warn("The 'normed' kwarg is deprecated, and has been "
C:\ProgramData\Anaconda3\lib\site-packages\matplotlib\axes\_axes.py:6462: UserWarning: The 'normed' kwarg is deprecated, and has been "
warnings.warn("The 'normed' kwarg is deprecated, and has been "
```



1.2.2 2.2 Word Cloud

```
In [10]: ham_words = ''
spam_words = ''
spam = data[data.label == 1]
ham = data[data.label ==0]

In [11]: import nltk

In [12]: for val in spam.text:
text = val.lower()
tokens = nltk.word_tokenize(text)
for words in tokens:
spam_words = spam_words + words + ' '


for val in ham.text:
text = val.lower()
tokens = nltk.word_tokenize(text)
for words in tokens:
ham_words = ham_words + words + ' '

In [13]: from wordcloud import WordCloud

In [14]: spam_wordcloud = WordCloud(width=600, height=400).generate(spam_words)
ham_wordcloud = WordCloud(width=600, height=400).generate(ham_words)

In [15]: #Spam Word cloud
plt.figure( figsize=(10,8), facecolor='k')
plt.imshow(spam_wordcloud)
```

```
plt.axis("off")
plt.tight_layout(pad=0)
plt.show()
```



```
In [16]: #legitimate/ham word cloud
plt.figure( figsize=(10,8), facecolor='k')
plt.imshow(ham_wordcloud)
plt.axis("off")
plt.tight_layout(pad=0)
plt.show()
```



```
In [19]: corpus[0]
```

```
Out[19]: 'Go jurong point crazi avail bugi n great world la e buffet cine got amor wat'
```

1.4 4. Create a sparse matrix by vectorizing words

```
In [20]: from sklearn.feature_extraction.text import CountVectorizer
```

```
vectorizer = CountVectorizer()
```

```
In [21]: X, y = vectorizer.fit_transform(corpus).toarray(), data['label'].values
```

```
In [22]: X[1]
```

```
Out[22]: array([0, 0, 0, ..., 0, 0, 0], dtype=int64)
```

1.5 5. Split the data

```
In [23]: # Splitting the dataset into the Training set and Test set
```

```
from sklearn.model_selection import train_test_split
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
```

```
In [24]: print('Number of rows in the total set: {}'.format(data.shape[0]))
print('Number of rows in the training set: {}'.format(X_train.shape[0]))
print('Number of rows in the test set: {}'.format(X_test.shape[0]))
```

```
Number of rows in the total set: 5572
```

```
Number of rows in the training set: 3900
```

```
Number of rows in the test set: 1672
```

1.6 6. Bayes with MultinomialNB

```
In [25]: from sklearn import tree
from sklearn import model_selection
from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score
from sklearn.tree import DecisionTreeClassifier
```

```
In [26]: from sklearn.naive_bayes import MultinomialNB
classifier = MultinomialNB()
```

```
classifier.fit(X_train, y_train)
predictions = classifier.predict(X_test)
print(predictions)
```

```
[0 0 0 ... 0 0 0]
```

```
In [27]: # Validate the accuracy of the predictions
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
print('Accuracy score: ', format(accuracy_score(y_test, predictions)))
print('Precision score: ', format(precision_score(y_test, predictions)))
print('Recall score: ', format(recall_score(y_test, predictions)))
print('F1 score: ', format(f1_score(y_test, predictions)))
```

```
Accuracy score: 0.9766746411483254
Precision score: 0.8874458874458875
Recall score: 0.9403669724770642
F1 score: 0.9131403118040089
```


We observed the predictions to be quite accurate even when the dataset was split into halves for train and test.

1.6.1 6.1 Confusion matrix

```
In [28]: conf_mat = confusion_matrix(y_test, predictions)
conf_mat_normalized = conf_mat.astype('float') / conf_mat.sum(axis=1)[:, np.newaxis]

sns.heatmap(conf_mat_normalized)
plt.ylabel('True label')
plt.xlabel('Predicted label')

Out[28]: Text(0.5,15,'Predicted label')
```



1.6.2 6.2 Using the model in new set of data

```
In [29]: data_new = pd.read_csv("./data/spam_source_2.csv",encoding='latin-1')
```

```
In [30]: data_new.head(5)
```

```
Out[30]:
```

	type	text
0	ham	Hope you are having a good week. Just checking in
1	ham	K..give back my thanks.
2	ham	Am also doing in cbe only. But have to pay.
3	spam	complimentary 4 STAR Ibiza Holiday or Â£10,000...
4	spam	okmail: Dear Dave this is your final notice to...

```
In [31]: # select and rename the columns relevant to us
data_new = data_new.rename(columns={"type":"label"})

print(data_new.groupby('label').size())
```

```
label
ham      4812
spam      747
dtype: int64
```

```
In [32]: data_new.isnull().values.any()
```

```
Out[32]: False
```

```
In [33]: data_new.head(5)
```

```
Out[33]:
```

	label	text
0	ham	Hope you are having a good week. Just checking in
1	ham	K..give back my thanks.
2	ham	Am also doing in cbe only. But have to pay.
3	spam	complimentary 4 STAR Ibiza Holiday or Â£10,000...
4	spam	okmail: Dear Dave this is your final notice to...

```
In [34]: # convert the output data/ label to binary value
data_new['label'] = data_new.label.replace(to_replace=['ham', 'spam'], value=[0, 1])
```

```
In [35]: corpus_new = []
for i in range(0,len(data_new)):
    mess_new = re.sub('[^a-zA-Z]',repl = ' ',string = data_new['text'][i])
    mess_new.lower()
    mess_new = mess_new.split()
    mess_new = [portstemmer.stem(word) for word in mess_new if word not in set(stop)]
    mess_new = ' '.join(mess_new)
    corpus_new.append(mess_new)
```

```
In [36]: X_new, y_new = vectorizer.transform(corpus_new), data_new['label'].values
```



```

In [37]: X_new.shape
Out[37]: (5559, 6304)

In [38]: y_new.shape
Out[38]: (5559,)

In [39]: predictions_new = classifier.predict(X_new)
         print(predictions_new)

[0 0 0 ... 1 1 0]

In [40]: # Validate the accuracy of the predictions
         print('Accuracy score: ', format(accuracy_score(y_new, predictions_new)))
         print('Precision score: ', format(precision_score(y_new, predictions_new)))
         print('Recall score: ', format(recall_score(y_new, predictions_new)))
         print('F1 score: ', format(f1_score(y_new, predictions_new)))

Accuracy score:  0.9883072495053067
Precision score:  0.94750656167979
Recall score:    0.9665327978580991
F1 score:        0.9569251159708415

```

We observed the model to perform good predictions on the spam dataset from different source

1.6.3 6.2 Save the model

```

In [41]: import pickle

         pickle.dump(classifier, open("bayes_classifier.pickle", "wb"))

         pickle.dump(vectorizer, open("vector.pickle", "wb"))

         pickle.dump(stop, open("stop.pickle", "wb"))

In [42]: # test load the classifier
         classifier = pickle.load(open('bayes_classifier.pickle', 'rb'))

         vectorizer = pickle.load(open('vector.pickle', 'rb'))

         stop = pickle.load(open('stop.pickle', 'rb'))

         predictions = classifier.predict(X_test)
         print('Accuracy score: ', format(accuracy_score(y_test, predictions)))

Accuracy score:  0.9766746411483254

```

1.6.4 Function to corpus a input message

```
In [43]: def corpus_message(text):
        mess = re.sub('[^a-zA-Z]', repl = ' ', string = text)
        mess.lower()
        mess = mess.split()
        mess = [portstemmer.stem(word) for word in mess if word not in set(stop)]
        mess = ' '.join(mess)
        mess_array= np.array([mess])
        print(mess_array[0])
        return mess_array
```

1.6.5 Custom text test

```
In [44]: test_string = 'text for free prize'
        test_corpus = corpus_message(test_string)
        test_vector = vectorizer.transform(test_corpus)
        test_vector
```

text free prize

```
Out[44]: <1x6304 sparse matrix of type '<class 'numpy.int64'>'
        with 3 stored elements in Compressed Sparse Row format>
```

```
In [45]: predictions = classifier.predict(test_vector)
        predictions
```

```
Out[45]: array([1], dtype=int64)
```

1.7 7. Ensemble Random Forest classifier

```
In [46]: # create and train the model
```

```
from sklearn.ensemble import RandomForestClassifier
model = RandomForestClassifier()
model.fit(X_train,y_train)
```

```
Out[46]: RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
        max_depth=None, max_features='auto', max_leaf_nodes=None,
        min_impurity_decrease=0.0, min_impurity_split=None,
        min_samples_leaf=1, min_samples_split=2,
        min_weight_fraction_leaf=0.0, n_estimators=10, n_jobs=1,
        oob_score=False, random_state=None, verbose=0,
        warm_start=False)
```

```
In [47]: predictions = model.predict(X_test)
        predictions
```

```
Out[47]: array([0, 0, 0, ..., 0, 0, 0], dtype=int64)
```


```
In [48]: # Validate the accuracy of the predictions
print('Accuracy score: ', format(accuracy_score(y_test, predictions)))
print('Precision score: ', format(precision_score(y_test, predictions)))
print('Recall score: ', format(recall_score(y_test, predictions)))
print('F1 score: ', format(f1_score(y_test, predictions)))
```

```
Accuracy score:  0.965311004784689
Precision score:  0.9878048780487805
Recall score:    0.7431192660550459
F1 score:        0.8481675392670157
```

```
In [49]: conf_mat = confusion_matrix(y_test, predictions)
conf_mat_normalized = conf_mat.astype('float') / conf_mat.sum(axis=1)[:, np.newaxis]

sns.heatmap(conf_mat_normalized)
plt.ylabel('True label')
plt.xlabel('Predicted label')
```

```
Out[49]: Text(0.5,15,'Predicted label')
```



1.8 8. Web App

```
In [50]: %%writefile mywebapp.py
```

```

from flask import Flask, render_template, request
from wtforms import Form, TextAreaField, validators
import pickle
import os
import re
import numpy as np
from nltk.stem.porter import PorterStemmer

portstemmer = PorterStemmer()

# loading the serialized classifier, vectorizer and stopwords
classifier = pickle.load(open('bayes_classifier.pickle', 'rb'))

vectorizer = pickle.load(open('vector.pickle', 'rb'))

stop = pickle.load(open('stop.pickle', 'rb'))

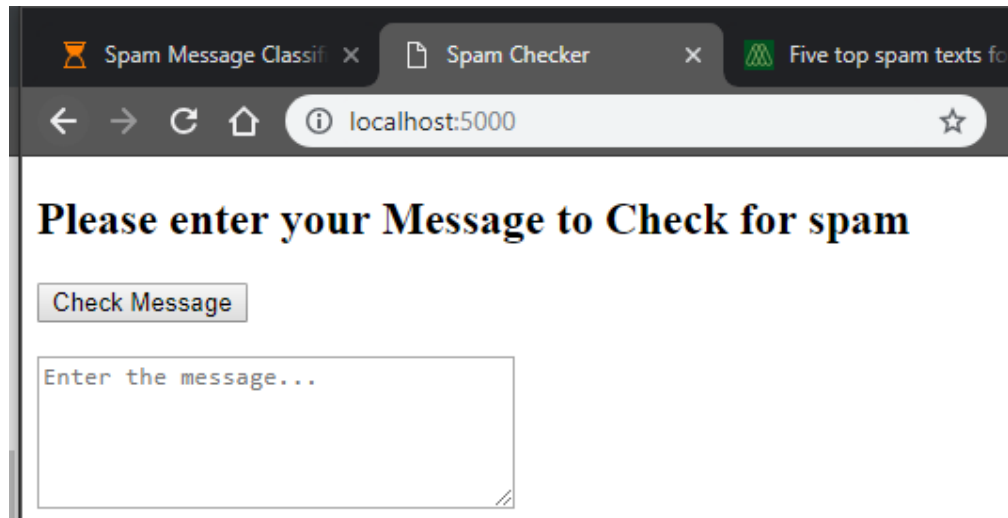
def corpus_message(text):
    mess = re.sub('[^a-zA-Z]', repl = ' ', string = text)
    mess.lower()
    mess = mess.split()
    mess = [portstemmer.stem(word) for word in mess if word not in set(stop)]
    mess = ' '.join(mess)
    mess_array= np.array([mess])
    return mess_array

def classify(message):
    label = {0: 'ham', 1: 'spam'}
    document = corpus_message(message)
    X = vectorizer.transform(document)
    y = classifier.predict(X)[0]
    proba = np.max(classifier.predict_proba(X))
    return label[y], proba

# create an instance (our app)
app = Flask(__name__)

@app.route('/', methods=['GET', 'POST'])
def index():
    form = None
    if request.method == 'POST' and 'spam' in request.form:
        form = request.form['spam']
    return render_template('default.html', form=form)

```



```
@app.route('/results', methods=['POST'])
def results():
    form = request.form
    if request.method == 'POST':
        message = request.form['spam']
        y, proba = classify(message)
        return render_template('results.html',
                               content=message,
                               prediction=y,
                               probability=round(proba*100, 2))
    return render_template('results.html', name=name)

if __name__ == '__main__':
    app.run(port=5000, debug=True);
```


Overwriting mywebapp.py

```
In [51]: #!/python mywebapp.py
```

1.9 9. Conclusion

The predictions generated from both the models were quite remarkable on both the test and external dataset. However, the dataset was not overly complex to warrant huge failures in predictions. Finally, we are happy with the results obtained.

1.10 10 Web App preview




Your message predicted for spam:

IMPORTANT - You could be entitled up to £3,160 in compensation from mis-sold PPI on a credit card or loan. Please reply PPI for info or STOP to opt out.

Prediction:

This message is **spam** (probability: 99.99%).

[Check another message](#)



Your message predicted for spam:

complete this survey to win a prize

Prediction:

This message is **spam** (probability: 99.31%).

[Check another message](#)

Your message predicted for spam:

call me

Prediction:

This message is **ham** (probability: 67.49%).

[Check another message](#)